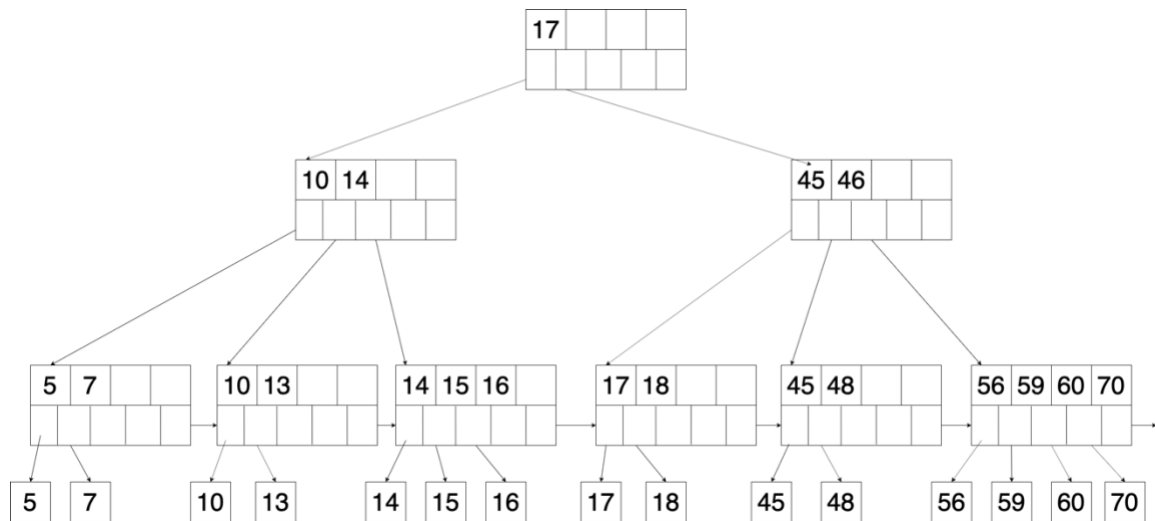


1. (a).

- Start at the root, read the root.
- Proceed down, to the leaf.
- Find the first leaf in the range and read the leaf with node 10 and node 13.
- Continue to read next leaf with node 17 in next block and read the leaf with node 18.
- Continue to read next leaf with node 45 in next block and read the leaf with node 48.
- Since next node value is 56, which is larger than 50, stop it.

There are 4 blocks I/O's are needed for the process.

(b).



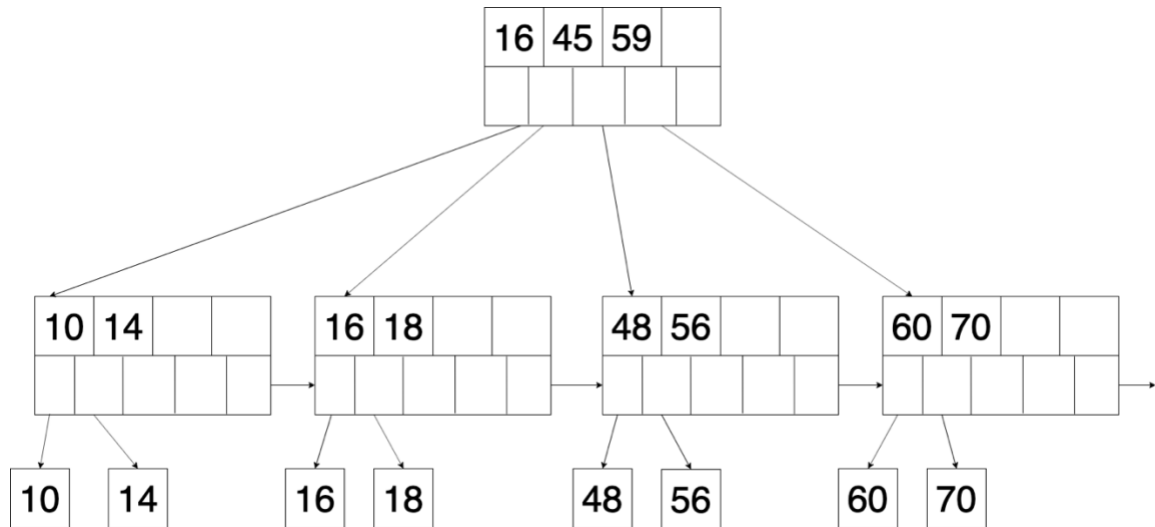
There are 3 blocks I/O's are needed for inserting 14. (read 2 blocks and write 1 block)

There are 3 blocks I/O's are needed for inserting 16. (read 2 blocks and write 1 block)

There are 7 blocks I/O's are needed for inserting 15. (read 2 blocks and write 5 block)

So total cost is 13 blocks.

(c).



There are 7 blocks I/O's are needed for deleting 5. (read 5 blocks and write 2 block)

There are 3 blocks I/O's are needed for deleting 7. (read 2 blocks and write 1 block)

There are 6 blocks I/O's are needed for deleting 13. (read 3 blocks and write 3 block)

There are 6 blocks I/O's are needed for deleting 15. (read 4 blocks and write 2 block)

There are 6 blocks I/O's are needed for deleting 17. (read 3 blocks and write 3 block)

There are 7 blocks I/O's are needed for deleting 45. (read 4 blocks and write 3 block)

There are 3 blocks I/O's are needed for deleting 59. (read 2 blocks and write 1 block)

So total cost is 38 blocks.

2. (a). for each (102-2) blocks b_r of R do

for each blocks b_s of S do

for each tuple r in b_r do

for each tuple s in b_s do

if r and s join then output (r, s)

Read R once: cost $B(R)$

Outer loop runs $B(R)/(M-2)$ times,

and each time need to read S: costs $B(R)B(S)/(M-2)$

Cost: $B(R) + B(R)B(S)/(M-2) = 20000 + 10000 * 20000 / (102 - 2) = 2020000$ blocks I/O.

(b). for each (102-2) blocks b_s of S do

for each blocks b_r of R do

for each tuple s in b_s do

for each tuple r in b_r do

if s and r join then output (s, r)

Read S once: cost $B(S)$

Outer loop runs $B(S)/(M-2)$ times,

and each time need to read R: costs $B(S)B(R)/(M-2)$

Cost: $B(S) + B(S)B(R)/(M-2) = 10000 + 20000 * 10000 / (102 - 2) = 2010000$ blocks I/O.

- (c). -Sort R into $B(R) / 100 = 20000 / 100 = 200$ runs. Cost: $2B(R)$
 -Sort S into $B(S) / 100 = 10000 / 100 = 100$ runs. Cost: $2B(S)$
 -Further merge R into $200 / 100 = 2$ runs. Cost: $2B(R)$
 -Further merge S into $100 / 100 = 1$ runs. Cost: $2B(S)$
 -Finally join 2 runs from R and 1 run from S. Cost: $B(R) + B(S)$
 Total Cost: $5B(R) + 5B(S) = 5 * 20000 + 5 * 10000 = 150000$ blocks I/O.
- (d). -Sort R into $B(R) / 100 = 20000 / 100 = 200$ runs. Cost: $2B(R)$
 -Sort S into $B(S) / 100 = 10000 / 100 = 100$ runs. Cost: $2B(S)$
 -Further merge R into $200 / 100 = 2$ runs. Cost: $2B(R)$
 -Further merge R into 1 run. Cost: $2B(R)$
 -Further merge S into $100 / 100 = 1$ runs. Cost: $2B(S)$
 -Finally join 1 run from R and 1 run from S. Cost: $B(R) + B(S)$
 Total Cost: $7B(R) + 5B(S) = 7 * 20000 + 5 * 10000 = 190000$ blocks I/O.
- (e). Hash table S into $(102 - 1) = 101$ buckets, send all buckets to disk.
 Hash table R into $(102 - 1) = 101$ buckets, send all buckets to disk.
 Join every pair of 101 buckets from table S and 101 buckets from table R.
 Cost: $3B(R) + 3B(S) = 3 * 20000 + 3 * 10000 = 90,000$ blocks I/O.
- (f). for each $(102-2)$ blocks b_r of R do
 for each tuple r_r of b_r do
 fetch corresponding tuple r_s of S
 Cost: $B(R) + T(R)B(S) / V(S, a) = 20000 + 200000 * 10000 / 100 = 20020000$ blocks I/O

Partitioned-hash join algorithm is most efficient.